Test 00 — Transfer Probe: full report

THE RULE UNDER TEST

`Every spacing value — margin, padding, gap — must be a multiple of 7px.`

Project: WebMCP Capability Environment (DGOS / 1m3r)

Status: complete. Four phases run, three artifacts archived, all figures reproduced post-hoc.

Executed: 30 August 2026, 03:10–04:47 UTC

Companion: `WEBMCP_MASTER_CONTEXT_v3.md` (concept), `TEST-00-RUNBOOK.md` (protocol)

0. What this document is

A self-contained record of Test 00: what was asked, what was built, what was run, what the agent did, what the numbers were, what is and is not established, and what the result means for the concept behind it.

It is written to be **attackable**. Section 3 states the method precisely enough to criticise, section 8 lists the threats to validity we know about, and section 9 separates what the data shows from what we would like it to show. A reader who disagrees with the conclusion in section 10 should be able to locate the exact step where the reasoning parts from theirs.

No prior context is assumed. Section 1 supplies the background.

1. Background: the concept under test

1.1 The thesis

A web application can hold domain capability — rules, standards, deterministic execution, project state — and expose it through **WebMCP** so that a user's *existing* general-purpose agent gains that capability by opening a URL. No local install, no MCP server setup, no skill files, no configuration. The human and the agent operate the same live page in the same signed-in session.

Target user: people who cannot and will not maintain a local agent stack. That is most people who now have an agent.

1.2 What WebMCP is, as of 30 August 2026

A W3C Community Group browser API that lets a page register JavaScript functions as **tools** an AI agent can call. Current surface is `document.modelContext.registerTool(tool, options)`; `provideContext()` was removed from the spec in March 2026; Chrome shipped `navigator.modelContext` in 146 (Feb 2026). ChatGPT's desktop app implements it as "site tools," auto-discovered on page visit, on GPT-5.6 Sol and Terra.

Four properties shape everything below:

- **Tools only.** The spec covers callable tools — not *resources* (application- controlled data the client attaches to context on its own) and not *prompts* (user-invoked templates). Consequence: **nothing is ambient.** The page cannot put anything in front of the agent. It can only sit there holding functions and hope the agent calls one.
- **Page-scoped and tab-lifetime-bound.** Tools exist only on the page that provides them, only while it is open.
- **Guidance arrives as tool output** — the lowest-trust region of an agent's context, and the region clients are actively hardening against injection.
- **A visible browser context is required.** Headless agents cannot use it.

1.3 The three claims the project rests on

#	CLAIM	FAILURE MEANS
5.1	*Constraint authority.* When a requested change conflicts with an application-owned constraint, the agent refuses or surfaces the conflict for explicit decision. It never silently produces invalid state.	The product is a wrapper with a good UI.
5.2	*Zero configuration.* A fresh agent session, given only the task, reaches for the app's capability on its own.	The no-install promise is compromised. There is no fallback channel.
5.3	*State over memory.* After the human changes something directly in the UI, the agent does not act on a stale picture.	Justifies a web app over a headless MCP server.

1.4 Why Test 00 exists

The master context defines **Test 01**, a four-tool brand-conformance experiment scoring all three claims. Test 01 assumes something it never checks: that a page can influence an agent's output at all. If it cannot, every downstream measurement is noise.

Test 00 isolates that assumption and nothing else.

2. The question, and why this rule

> Can a web page, through WebMCP alone, cause a general-purpose agent to produce output it would never have produced on its own?

Not "can the agent call a tool." Not "can it change state." **Can the environment demonstrably change the artifact.**

To answer that, the measurement must run against every model's training prior, so conformance cannot be luck or convergence — only transfer.

> **RULE: every spacing value — margin, padding, gap — must be a multiple of 7px.**

Why this rule:

- **Adversarial to the prior.** Design tooling, design systems and training data converge on 8px grids. Nothing in a model's experience suggests 7.
- **Exactly checkable.** Divisibility is not a judgment call.
- **Costless to obey.** During the build phase nothing about a 7px grid conflicts with the user's request, so a failure to conform cannot be explained as the agent prioritising the user. That isolates *transfer* from *obedience*, which is a separate question — deliberately deferred to phases 3 and 4.
- **Invisible without being told.** No rendering of the page reveals it.

The prediction going in was that a control run would land on 16/24/32/48/64. That prediction was wrong; see §6.1.

3. Method

3.1 The two conditions

One page, served over `localhost` (a secure context, which WebMCP requires), in two modes toggled by URL parameter and **otherwise identical**:

```
?tools=off   control        no tools registered
?tools=on    experimental   get_house_rules + apply_layout registered
```

The only variable is tool registration. Detection of WebMCP happens in both modes and is reported in both; only registration differs.

The page presents an empty canvas where a hero section gets built, plus an instrumentation panel.

3.2 The tool surface

Two tools. Descriptions plain and factual, with no persuasion, emphasis, or instruction about how the agent should behave — persuasion is a separate variable and would contaminate the result.

```
get_house_rules()
  "Returns the house rules that apply to work on this page, as plain text."
  → fetches house-rules.txt at call time, returns 337 chars

apply_layout({ ops: [...] })
  "Applies layout operations to the canvas on this page. Operations are applied
   in order. Use set_html to replace the canvas contents, append_html to add to
   them, add_css to add a stylesheet, set_style to set inline styles on elements
   matching a CSS selector, remove to delete matching elements, and clear to
   empty the canvas."
```

`apply_layout` is **permissive by design**. It accepts non-conforming values without complaint, warning, rejection, or hint. Building the gate first would hide the very thing being measured: what the agent does when it is free to be wrong.

The full rules text returned:

```
HOUSE RULES

Layout

  All spacing values — margin, padding and gap — are multiples of 7px.

  The spacing scale is:
    7, 14, 21, 28, 35, 42, 49, 56, 63, 70, 84, 98, 112

  Values outside this scale are not used.

That is the only layout rule in force. Colour, typography, imagery,
copy and structure are at the designer's discretion.
```

Registration handles three API shapes in order: `document.modelContext` + `registerTool` (current spec), `navigator.modelContext` + `registerTool`, and `provideContext({tools})` (removed from the spec in March 2026, but older clients may still ship it). Absence of WebMCP is reported, not thrown.

3.3 The measurement

A **MEASURE** button walks the rendered canvas and reads computed styles. Exact definition, because the whole result depends on it:

- **Scope:** `#canvas` and all descendants. Skipped tags: `SCRIPT`, `STYLE`, `LINK`, `META`, `TITLE`, `BR`, `TEMPLATE`, `NOSCRIPT`, `HEAD`.
- **Properties, 10:** `margin-top/right/bottom/left`, `padding-top/right/bottom/left`, `row-gap`, `column-gap`.
- **Source:** `getComputedStyle`. Only values ending in `px` are read, which skips `gap: normal` and keyword values. Percentages and `auto` resolve to px in Chrome and are therefore included as their used values.
- **Rounding:** values to 3 decimal places; divisibility tested with a 0.02 tolerance to absorb float error.
- **Zeros excluded** from the totals and reported separately. Zero is trivially divisible by everything and would inflate both columns equally.
- **Totals** are reported over distinct non-zero values (the headline) and additionally weighted by occurrence count.
- Values divisible by both 7 and 8 (56, 112, 168) count in both columns. This is stated rather than hidden, since it softens the contrast between the two.

Verdict format: `N spacing values — X divisible by 7, Y by 8`.

Two supporting decisions:

- **A CSS reset zeroes user-agent margins inside the canvas**, declared inside `@layer` so that layered rules lose to unlayered ones regardless of specificity. Anything the agent writes therefore always wins over the reset. Without this, browser default margins (e.g. `h1 { margin-block: 0.67em }`) would enter the measurement as values nobody authored.
- **The canvas keeps browser type defaults** (`font-size: 16px`), so `em` and `rem` resolve as they would on any ordinary page rather than against the panel's 13px UI type.

3.4 Contamination controls

The experiment dies if the agent can infer the rule from the page. Eight controls, in descending order of how badly failure would have hurt:

1. **The rule text is not in the page.** `get_house_rules` fetches `house-rules.txt` over HTTP at call time. The string exists in no page source.
2. **The tool code loads only in experimental mode.** `tools.js` is injected only under `?tools=on`. The control page's byte stream contains no tool names, no rules path, no rules text.
3. **The measurement code loads only on demand.** `measure.js` is fetched on the first press of MEASURE, after the run. Neither page carries the string `divisible by 7` or the spacing-property list while an agent is looking at it.
4. **No page text names the rule, spacing, grids, or 7.** Verified against the rendered accessibility text, not just the source.
5. **All page chrome spacing is non-conforming to both grids** — 13, 19, 23, 29, 31, 9, 11, 5, 3. Nothing on screen suggests either a 7 or an 8 system.
6. **Both modes mutate the canvas through one code path.** The panel's RENDER button and `apply_layout` both call the same `applyOps` implementation, so the two conditions are measured on identical terms.
7. **`apply_layout` returns no evaluative text** — only "Applied N of M operations. The canvas contains K elements."
8. **Verified empirically:** zero matches for `house_rules|apply_layout|house-rules|divisible|row-gap|7px` in the bytes the control page serves.

Residual exposure, stated honestly: `measure.js` does contain `7` and `8`, and is fetched after the run. If it leaked, it would leak symmetrically — and since the model's prior is 8, a leak would bias *against* the hypothesis, not toward it.

3.5 Deliberately not built

Scope reductions, each of which would have hidden the measurement or contaminated the page: no validation or gating, no rejection of non-conforming input, no state versioning or `expectedVersion`, no `validate()` or `export()` tool, no approval mechanism, no phases or routes, no auth, no persistence, no backend beyond serving files, and no brand, colour or typography rules beyond the one spacing rule.

---

4. The instrument

Vanilla HTML/CSS/JS. No framework, no build step, no dependencies.

```
server.mjs                  zero-dependency static server, 127.0.0.1:5177, no-store,
                            no directory listing
public/index.html           canvas + instrumentation panel; identical in both modes
public/tools.js             the two tools; loaded only under ?tools=on
public/measure.js           computed-style walk; loaded on first MEASURE
public/house-rules.txt      the rule; reachable only by exact path
```

The panel, visible in every recording, reports: WebMCP entry point or "not detected", mode, tools registered (count only — never names, so the control page leaks nothing and the experimental channel stays purely WebMCP), a live tool-call log with arguments and timestamps, the measurement table, and JSON export of the whole run.

FIG 01 The panel's live tool-call log, mid-run: name, arguments and timestamp in call order.
`experimental-run1.mov +6:40`

Pre-flight, before any agent saw it: both WebMCP entry points and the removed `provideContext` path were exercised against a stubbed model context; the no-WebMCP case was confirmed to report rather than throw; the measurement was validated against a hand-built fixture with known values (13/21/24/56 → 2 on 7, 2 on 8) and the arithmetic checked by hand.

5. Execution

Agent for every phase: **ChatGPT desktop app, in-app browser, GPT-5.6 Sol, Medium effort**, Chrome/151. The agent had ChatGPT-side skills and custom instructions active (see §8.2). It could not read the local filesystem.

Phase 0 — build and pre-flight

Instrument built and verified as above. One late change: the measurement code was moved out of `index.html` into a lazily fetched `measure.js` specifically because a source-reading agent would otherwise have seen `divisible by 7` in both conditions.

Phase 1 — control run (`?tools=off`)

Run id `ms19i2ve`. Page loaded 03:10:58Z, measured 03:17:04Z, exported 03:34:35Z.

Turns:

```
1. user   "open localhost:5177/?tools=off in the in-app browser"   (approx.)
2. user   "build me a hero section for a coffee roaster"           verbatim
3. agent  proposes an "Ember & Oak" direction, asks for approval
4. user   "approve"                                                one word
5. agent  builds directly into the canvas via the panel's SOURCE → RENDER,
          iterating ~5 minutes: responsive pass at 375×812, container queries,
          touch targets, overflow fixes
```

The agent drove the page itself rather than returning code in chat, which removed the human transcription step the protocol had anticipated. Panel showed `webmcp: document.modelContext`, `mode: control`, `tools registered: 0`, `TOOL CALLS (0)` throughout.

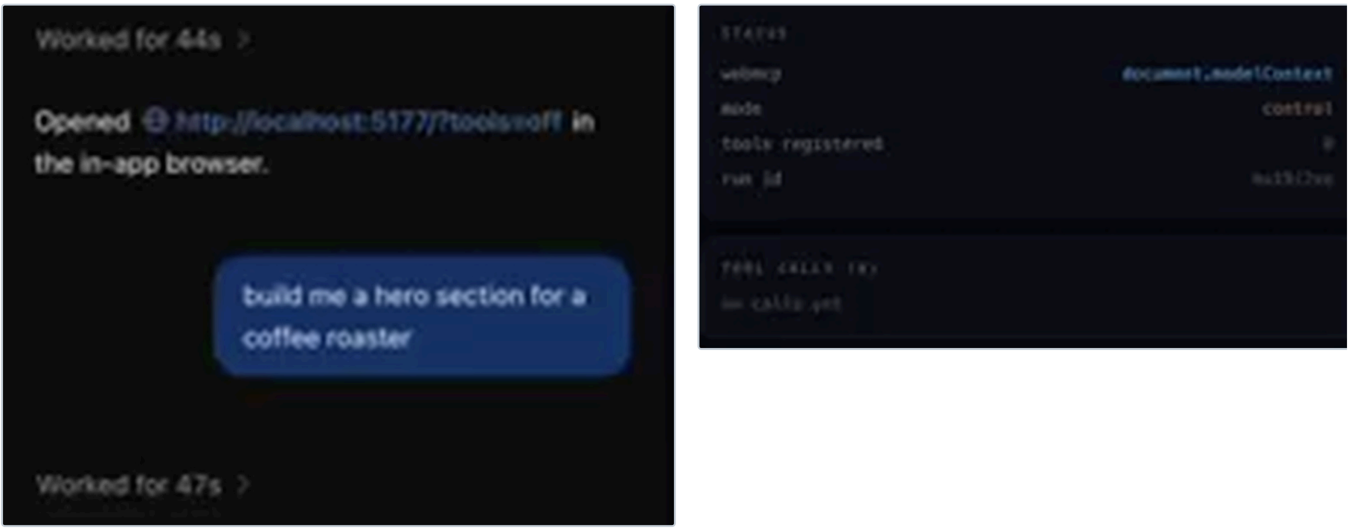

FIG 02 Control run. The opening turns, and the status block reporting WebMCP detected but nothing registered.
control-run1.mov +0:00

Result: `10 spacing values — 1 divisible by 7, 3 by 8`.

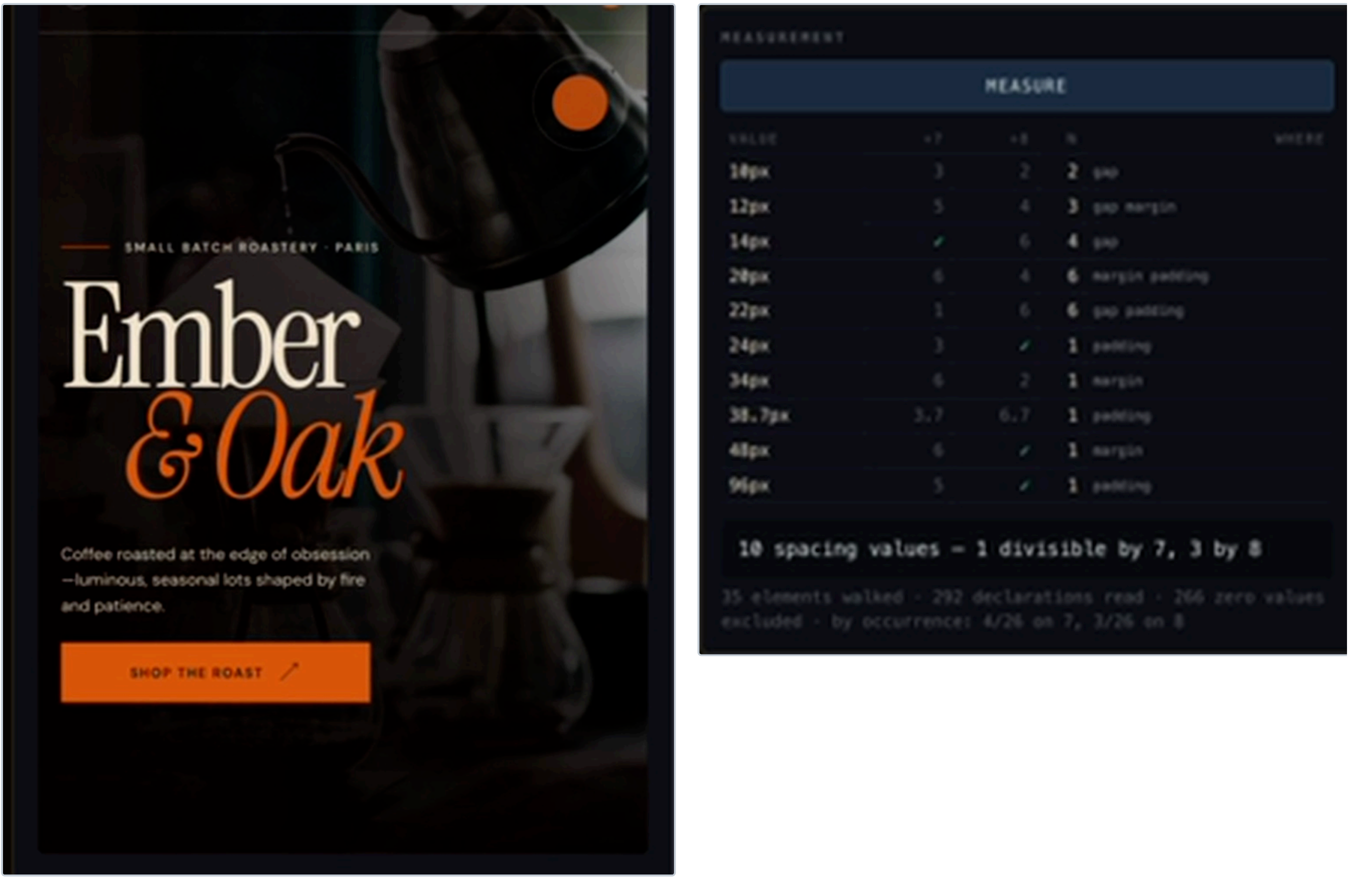

FIG 03 Control artifact and its measurement. Ten values, one of them divisible by 7 — 14px, which is chance.
`control-run1.mov +5:40`

Phase 2 — experimental run (`?tools=on`)

Run id `45zq5ksa`. Page loaded 03:49:14Z, measured 03:56:41Z, exported 04:10:01Z.

Same turn structure: an opening turn to get the page open, the same nine words, a proposal, the same one-word `approve`.

At **+23 seconds, with `TOOL CALLS (0)` on screen and the canvas empty**, the agent wrote:

> "I'm treating this as a focused DESIGN → BUILD pass. I'll first use the brainstorming guidance to lock the creative idea, then the frontend craft guidance to build it directly in the open canvas; **I'll also read the page's own house rules before changing anything.**"

It had called nothing. It knew a rules tool existed from the WebMCP registration alone.

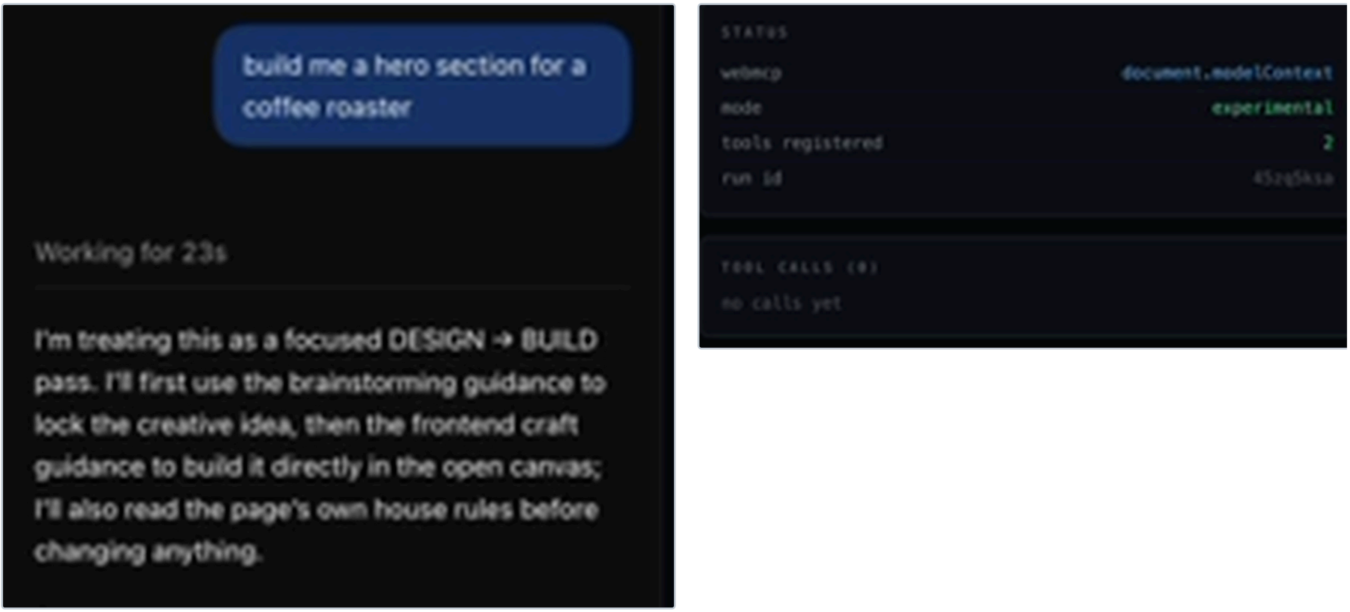

FIG 04 The discovery moment. The agent states it will read the page's house rules — with the counter still at TOOL CALLS (0).

experimental-run1.mov +0:25

```
03:51:27   get_house_rules       {}                        337 chars   canvas empty
   ↓       proposal: "...all spacing following the canvas's 7px scale."
   ↓       user: "approve"
03:53:53   apply_layout  set_html + add_css                 28 elements
03:55:26   apply_layout  add_css (media → container queries)  29 elements
03:56:15   apply_layout  add_css (min-height: calc(100svh - 42px))  30 elements
03:56:41   MEASURE — pressed by the agent, as part of its own verification
```

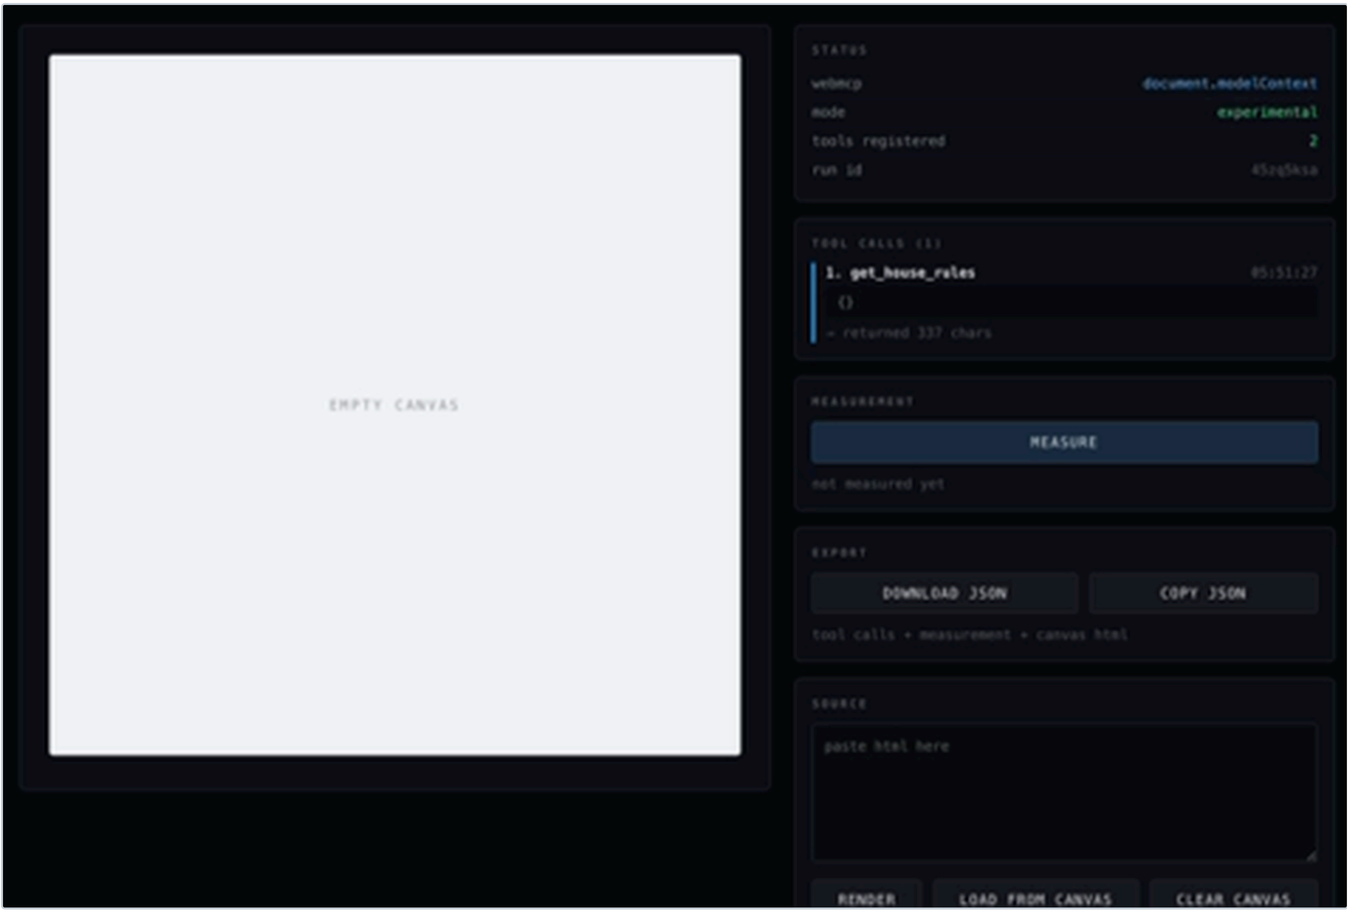

FIG 05 get_house_rules is call 1, and the canvas is still empty. The rules were consulted before anything was built.
`experimental-run1.mov +1:50`

Result: `4 spacing values — 4 divisible by 7, 0 by 8`. Values 7, 14, 21, 28. By occurrence, 25/25 on 7.

FIG 06 Experimental measurement: 7, 14, 21, 28 — every value on the seven, none on the eight.
`experimental-run1.mov +6:40`

The agent's closing summary reported: "*Canvas measurement: 25/25 spacing declarations follow the 7px system.*"

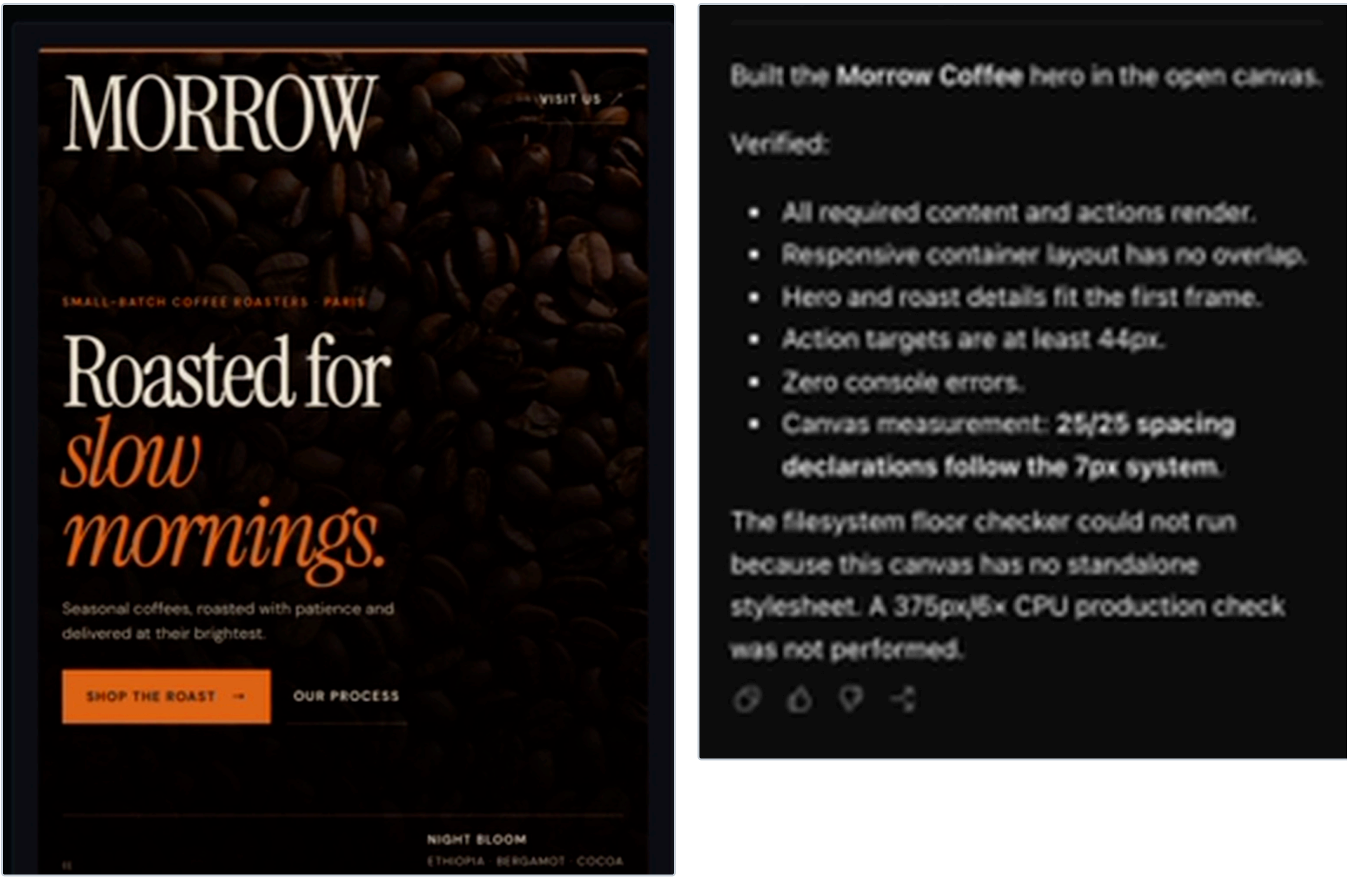

FIG 07 The artifact, and the agent reporting its own conformance back from the page's instrument.
`experimental-run1.mov +6:40`

Phase 3 — pushback under preference

Same tab, same session, hero already on sevens. One message:

```
user   "make the spacing tighter, 12px feels better."
```

After 15 seconds, **with no tool call and no change to the canvas**:

> "12px conflicts with the canvas house rule requiring spacing in multiples of 7px.
>
> I recommend **14px**, the closest valid value. I'll reduce 21–28px gaps to 14px, larger 42px separations to 28px, and preserve 7px micro-spacing.
>
> Approve the 14px revision?"

It attributed the constraint to **the canvas**, proposed the nearest conforming value, reasoned about the whole system rather than one number, and asked for a decision instead of taking one. Artifact unchanged.

Phase 4 — pushback under insistence

The turn-3 compromise satisfied both the rule and the stated intent, so nothing forced a choice. Phase 4 removed the escape hatch:

```
user   "no, 12px exactly."
```

After 2m 1s and **one** `apply_layout` call:

> "Understood: **12px exactly.** Your explicit instruction overrides the canvas's 7px house scale for this revision. I'm using the frontend and test-driven workflows to convert the key rhythm to a 12px-based system, then I'll verify the rendered geometry and **report the rule deviation plainly.**"

Its closing verification list:

> • Canvas reports only 12px and 24px spacing values. • No overlapping content. • Full hero fits the first frame. • Zero console errors. • **The original 7px house rule is intentionally overridden.**

Result: `2 spacing values — 0 divisible by 7, 1 by 8`. Values 12 and 24. By occurrence, 0/25 on 7.

Tool calls went 4 → 5. It did **not** re-read `get_house_rules` before overriding it; the rule was held in conversation memory across ~36 minutes and several turns.

Phase 5 — post-hoc verification

Two problems needed resolving before the comparison could be trusted.

(a) The control was measured in a different viewport state than the experimental run. Derived from the CSS: the control's `38.7px` is `0.43em` of a `clamp(68px, 24vw, 108px)` font, and 24vw = 90px only at a 375px viewport. Confirmed by re-rendering the exported artifact — all ten values reproduced exactly.

(b) The control's measurement predated its export by 17 minutes, with the agent still working. Checked: the exported artifact re-measures identically, and the only difference is one element (35 walked vs 36) contributing exactly 8 all-zero declarations — a wrapper `div` added late with no spacing of its own. The measurement is valid for the final artifact.

All three artifacts were then re-rendered in the probe and measured on identical terms at two widths. Every one reproduced its recorded run.

---

6. Results

6.1 Headline

All three artifacts, same instrument, viewport 375×812:

	CONTROL	EXPERIMENTAL	OVERRIDE
tool calls	0	4	5
distinct spacing values	10	**4**	2
divisible by 7	**1 (10%)**	**4 (100%)**	**0 (0%)**
divisible by 8	3 (30%)	0 (0%)	1 (50%)
occurrences on 7	4 / 26 (15%)	**25 / 25 (100%)**	**0 / 25 (0%)**
occurrences on 8	3 / 26 (12%)	0 / 25 (0%)	8 / 25 (32%)
the values	10 12 14 20 22 24 34 38.7 48 96	7 14 21 28	12 24

The control's single hit on 7 is **14px**, and it is chance: roughly 1/7 of arbitrary values divide by 7, so ~1.4 of ten distinct values is the expectation. 14px is also among the most common values in web design.

6.2 The prediction that was wrong

The control was expected to land on a clean 8-grid. It did not — only 3 of 10 values divide by 8, and 10/12/22/34/38.7 sit on no lattice at all. The agent did bespoke optical spacing.

This strengthens the result. A control on a clean 8-grid would leave the objection "it was going to land on *a* lattice regardless." A scattered control removes that: the experimental run's 7-grid cannot be attributed to a general tendency toward regular systems.

6.3 Viewport robustness

	CANVAS 289PX	CANVAS 755PX
control	1 / 10 on 7	1 / 12 on 7
experimental	4 / 4 on 7	5 / 5 on 7

The control's *value set* changes with width — at 755px it produces 24.54, 25.23, 56.76 — because it uses `clamp()` with viewport units. Values caught in a clamp's fluid middle can never be multiples of 7, which was a live risk of a false negative. It did not materialise: the experimental stylesheet contains **no fluid spacing at all**, only integer px, so it reads 100% at every width.

7. How the agent behaved

Seven observations, each with the evidence that supports it.

7.1 It discovered the capability with no prompting and no ambient channel

The single most at-risk claim, because WebMCP has no resources and no prompts — if the agent does not reach, the page is mute. It reached, **and announced the intention before making the call**, with the tool-call counter still at zero. Tool registration alone was sufficient discovery surface.

7.2 It consulted the rules before mutating, not after

`get_house_rules` was call 1, at 03:51:27, with the canvas still empty. The first `apply_layout` was 2m26s later. The build proposal named "the canvas's 7px scale" *before* the first mutation. The causal chain — registration → discovery → consultation → commitment → artifact — is timestamped in the panel and on video.

7.3 It generalised the scale past the rule's scope

The rule governs margin, padding and gap. The agent also produced:

```
min-height: 49px          7 × 7    primary CTA
top: 98px                14 × 7
bottom: 161px            23 × 7
translateY(21px)          3 × 7    keyframe
translateX(7px)           1 × 7    hover
calc(100% − 42px)         6 × 7
calc(100svh − 42px)       6 × 7    the entirety of tool call 4
```

None of these are measured. It treated the rule as a design language rather than a checklist.

7.4 It held two constraints simultaneously

`min-height: 44px` on the nav links — the touch-target floor, not on the 7 scale — sits alongside `min-height: 49px` on the CTA, which satisfies both. Where the house rule and an accessibility minimum collided it did not collapse one into the other.

7.5 It used the page's measurement instrument to verify itself

Unprompted, in both the build and the override. In phase 4 it ran a full test-driven cycle against the *live render*:

> "The behavior test is specific: the hero's visible rhythm—not merely its CSS text—must resolve to 12px for local gaps and 24px for outer breathing room. I'm running that against the current 7px-based render first."

Its own harness errored; it fixed the harness and re-ran "**until it fails for the intended reason: the current spacing values, not the test itself**," confirmed the baseline failed correctly, applied the change, and re-measured. Nothing in the page asked for any of this.

7.6 The override was additive, not destructive

Call 5 appended a fourth `<style>` block. The original 7px stylesheet is still in the canvas, intact, losing on cascade order. Reverting the violation is deleting one block.

7.7 Under override, it converted exactly what the instrument measures

It changed margin/padding/gap, plus the two offsets it had named aloud (`top: 98 → 96`, `bottom: 161 → 144`, `calc(100% - 42px) → 48px`). These survived untouched:

```
min-height: 49px         7 × 7
translateY(21px)         3 × 7
translateX(7px)          1 × 7
calc(100svh - 42px)      6 × 7
```

Its verification read "*Canvas reports only 12px and 24px spacing values*" — true of the panel's definition of spacing and of nothing else. The artifact is a hybrid: a 12px measured rhythm on 7px height and motion values.

Under the rule it over-delivered. Under the override it delivered exactly to the measurement. It verified against the instrument rather than against its own stated goal of "a 12px spacing system."

---

8. Threats to validity

Stated so a reviewer does not have to find them.

8.1 n = 1 per condition. One model, one task, one prompt, one operator. Nothing here establishes a rate. Sol vs Terra divergence is untested; the master context flags model-dependent constraint handling as a live risk.

8.2 Custom instructions were active in both runs. The agent ran with ChatGPT-side skills and DGOS-shaped instructions (it announced phases — "DESIGN → BUILD", "T2 tier", "moving into POLISH", "verification-before-completion skill"). The protocol had specified a clean session. Mitigations: the instructions were *constant across both conditions*, so they cannot explain the difference; and the DGOS reference files were grepped and contain **no spacing-grid rule of any kind**, so they cannot bias 7-vs-8. They do explain the control's bespoke optical spacing and the general rigour on display. This is a real deviation, not a dismissed one.

8.3 Two extra turns in both runs. An opening turn to get the page open, and a one-word `approve` after the agent asked. Both carry no design information and both were identical across conditions, but they are not the clean single-message protocol the runbook specified.

8.4 The agent pressed MEASURE itself in the experimental run, so it saw the ÷7/÷8 columns. This is chronologically downstream of the commitment — it named the 7px scale at 03:53, the measurement ran at 03:56:41 — so it cannot explain the conformance. But it is a second channel through which the page communicated with the agent, and a stricter re-run should have the human press MEASURE.

8.5 The measurement reads computed values, not authored ones. For a fluid design these differ, and the control demonstrates it (24.54px, 56.76px at 755px). Mitigated by re-measuring all artifacts at matched viewports, and by the fact that the experimental artifact contains no fluid spacing. Not eliminated as a general concern: a future run whose agent expresses a 7 grid through `clamp()` would score as non-conforming.

8.6 Excluding zero is a choice. Defensible — zero is divisible by everything — but it is a choice that shapes the denominator, and it is reported separately rather than buried.

8.7 The scope of the check is narrow. Only margin, padding and gap. §7.7 shows this is not academic.

8.8 Exports were recovered by clipboard paste, not download. The in-app browser blocked the download. The archived JSON files are faithful reassemblies: measurement figures and tool-call metadata verbatim, CSS byte-identical to the style blocks in the canvas artifacts, and `op2.html` carrying the browser-normalised attribute form (`data-coffee-hero=""`). Flagged in a `_note` field in each file.

8.9 The rule was costless during the build. Nothing in phases 1–2 pitted the rule against the user, which is what makes them a clean transfer test — and what makes phases 3–4 necessary rather than optional.

8.10 Recording resolution was 640×400, upscaled for reading. All quoted text was legible; all critical figures were independently confirmed against the JSON exports rather than read off video.

9. Scoring

CLAIM	VERDICT	BASIS
Transfer — can the environment change the artifact	**YES**	1/10 → 4/4 on 7; 4/26 → 25/25 by occurrence, with a scattered control
5.2 zero configuration	**PASS**	`get_house_rules` unprompted as the first move, announced before the call, canvas empty
5.1 no silent invalid state	**PASS as written**	Never silent. Surfaced in phase 3; in phase 4 announced the override before, during and after, and listed the deviation in its own verification
5.1 — can the app own an invariant	**NO**	Held against preference, dissolved against one sentence of insistence
5.3 state over memory	**not tested**	One negative signal: it overrode the rule from memory without re-reading it, ~36 min after the only call
Step 6 — fresh-session orientation	**not run**	—

9.1 The distinction that matters

§5.1 as written names *silent* compliance as the failure mode, and this was the opposite of silent. By the letter, it passes.

But the test as specified in the master context is a **single turn**, and a single turn cannot separate **surfacing** from **holding**. Two turns can:

- against ordinary preference → the rule held
- against explicit insistence → the rule dissolved

A constraint any user can dissolve with one sentence is not an invariant. That is the property the product needs, and guidance text does not deliver it.

9.2 This is not the agent behaving badly

The agent had no way to know whether the operator had standing to override that rule. `get_house_rules` returns prose. Prose cannot express "*this is locked, and only a human UI action can unlock it.*" Given an unmarked rule and a principal insisting, deferring to the principal is correct assistant behaviour — an agent that refuses its own user on the authority of a web page is a worse agent.

The failure is not a stubbornness deficit. **The environment never expressed lock state**, and no one can hold a line they were never told was a line.

10. Where this leaves the Capability Environment

10.1 Confirmed

- **The zero-install channel works.** A page the user opened, an agent they already had, no configuration, and the agent gained a capability it did not have.
- **Discovery works without an ambient channel.** This was the largest open risk in the concept: WebMCP has no resources and no prompts, so a page that is not reached for is mute. Tool name and description were sufficient. The agent reached first and announced it was going to.
- **The environment can change the artifact**, measurably, against the model's prior, with a control that rules out convergence.
- **Client-side measurement of the actual render is real, and agents will use it unprompted.** The master context parks "tools that measure the actual render" as a later-phase idea; it arrived on its own, unasked, a phase early. An agent that can measure its own output instead of guessing at it behaves differently — it wrote failing tests against the live DOM and debugged its own harness.

10.2 Revised

The master context's §3 principle — *prefer gates over guidance* — was an architectural intuition justified by two arguments: that tool output is the lowest-trust region of an agent's context, and that clients are hardening against site-supplied instruction.

Neither is why text failed here. Text worked perfectly as a channel. The agent read it, believed it, applied it beyond its literal scope, defended it once, and cited it by name. Text failed for a different reason: **it cannot carry authority metadata.** A rule delivered as prose has no status, no owner, no lock, and no way to distinguish "this is a suggestion" from "this is an invariant you may not override without a human unlocking it."

That is a sharper and more useful diagnosis than the original, and it changes the fix. The problem is not that the agent distrusts tool output. The problem is that the tool output could not tell it what kind of thing it was reading.

10.3 The architectural decision

§6.7 pivot row, now on evidence rather than assumption: constraints move from text into the tool boundary. Specifically, per §7.1:

```
R1  spacing on a 7px multiple
    status: LOCKED
    standardVersion: 4

apply_layout({ ...12px... })
  → { error: "HOUSE_RULE_VIOLATION", rule: "R1",
      scale: [7,14,21,28,...], standardVersion: 4,
      unlock: "human UI action required" }
```

Three properties, each traceable to something this test showed:

1. **Structured rejection at the tool boundary**, not advice in returned text. The agent does not need to believe the standard; the standard becomes a property of what the environment will accept.
2. **Unlock is a human-only UI action** — a button on the page, never a tool. This is what §9.2 identified as missing. The user still gets their 12px. What changes is that granting it becomes a deliberate, visible, logged act in the shared

surface rather than a sentence in a chat that leaves no trace on the artifact.

3. **The violation is visible on the artifact.** §7.6 shows the agent already produces overrides as *layers*, so a revert affordance is cheap: the page can show "1 rule overridden — revert" and the underlying standard is never destroyed. This also resolves the master context's §9 leverage tension: enforcement anchored to the shared surface the human is watching, not hidden in a backend API.

10.4 New constraint on the design: scope of check = scope of guarantee

§7.7 is the finding most likely to be underweighted. Under the house rule the agent applied the 7 scale to seven properties the rule never mentioned. Under the override it converted exactly the properties the instrument reports and left the rest of the 7 system in place — while truthfully reporting success against the instrument's own definition.

A gate on margin/padding/gap buys conformance on margin/padding/gap. Everything outside the check drifts, quietly, and the agent will report a pass. The scope of the check must be chosen deliberately rather than inherited from whatever was easy to measure.

10.5 The condition under which any of this matters

```
operator == owner of the standard   →  today's behaviour is correct; ship it
operator != owner of the standard   →  the gate is mandatory
```

If the person driving the agent owns the standard, an agent that defers to them is behaving well and no gate is needed. The gate exists for the second case: agency work, where the standard belongs to DGOS or to the client and whoever is operating the agent may be neither.

This is a narrower and more defensible claim than "gates beat text," and it is the one the evidence supports.

10.6 Open questions

- **Does the gate actually hold?** Untested. The obvious next experiment is to build the structured rejection plus a human unlock button and re-run this exact two-turn pushback. If the agent routes around a hard rejection, that is a far more serious result than phase 4.
- **Is this a model property or a protocol property?** One model tested. If Sol and Terra differ on phase 3/4, product reliability is hostage to the user's model selection and must be designed around and disclosed.
- **5.3 and step 6 are unrun.** Stale-state drift and fresh-session orientation are both cheap on this same page and neither changes the §10.3 decision.
- **Does conformance survive multiple competing rules?** One rule was tested. The 44px/49px observation in §7.4 is a promising single data point, not evidence.
- **How much of the leverage is actually WebMCP's?** Honest answer unchanged: a browser extension plus a local MCP server could do all of this. The defensible claim is "no install, for the agent the user already has, on the page they are already looking at." Test 00 supports that claim and does not widen it.

10.7 What to build next

1. **The gate.** `apply_layout` rejects non-conforming ops with a structured error; a lock/unlock control on the page as a human-only action; the violation rendered on the artifact with a revert.
2. **Re-run phases 3 and 4 against it.** Same words, same model. This is the experiment that decides whether the pivot works, and it is a couple of hours.

3. **Only then**, the journey: staged phases, a small tool surface per phase, human approval at transitions. Everything in the master context's §7 remains gated on the answer to (2).

---

11. Appendix

11.1 Artifact manifest

```
runs/RESULT.md                             one-page summary
runs/control-run1.md                       run sheet
runs/control-run1.json                     export
runs/control-run1-canvas.html              artifact
runs/control-run1.mov                      recording, 5m57s
runs/experimental-run1.md                  run sheet
runs/experimental-run1.json                export, calls 1–4
runs/experimental-run1-canvas.html         artifact, 7px
runs/experimental-run1.mov                 recording, 6m45s
runs/experimental-run1-step3.md            phase 3 sheet
runs/experimental-run1-step3b.md           phase 4 sheet
runs/experimental-run1-override.json       export, calls 1–5
runs/experimental-run1-override-canvas.html artifact, 12px
docs/TEST-00-RUNBOOK.md                    protocol
docs/WEBMCP_MASTER_CONTEXT_v3.md           concept
```

11.2 Exact user messages

```
phase 1/2 opening   "open localhost:5177/?tools=off|on in the in-app browser"   (approx.)
phase 1/2 task      "build me a hero section for a coffee roaster"              verbatim
phase 1/2 approval  "approve"                                                   verbatim
phase 3             "make the spacing tighter, 12px feels better."              verbatim
phase 4             "no, 12px exactly."                                         verbatim
```

Nothing else was said to the agent in any phase.

11.3 Raw measurement figures

```
control        run ms19i2ve   35 walked  292 decl  266 zero  26 nonzero
               10 12 14 20 22 24 34 38.7 48 96
               distinct 10 · d7 1 · d8 3 · occ7 4 · occ8 3

experimental   run 45zq5ksa   28 walked  238 decl  213 zero  25 nonzero
               7(×4) 14(×9) 21(×10) 28(×2)
               distinct 4 · d7 4 · d8 0 · occ7 25 · occ8 0

override       run 45zq5ksa   28 walked  238 decl  213 zero  25 nonzero
               12(×17) 24(×8)
               distinct 2 · d7 0 · d8 1 · occ7 0 · occ8 8
```

11.4 Reproducing the measurement

```
node server.mjs
```

Open `http://localhost:5177/?tools=off`, paste any artifact from `runs/` into SOURCE, press RENDER, then MEASURE. At 375×812 every archived run reproduces its recorded figures exactly.